

PRODUCT OVERVIEW

The meeting is not the artefact.

Recap sits in a call, writes down what was said, tells you what to say next while it is still happening, and afterwards turns what you agreed to into a first step small enough to actually take.

Notes compete with the conversation

In any meeting that matters you are asked to do two incompatible things at once: take part, and keep a record. Whoever writes the best notes contributes the least, and whoever contributes most remembers the least. The record of what a team decided ends up depending on who happened to be typing.

Transcription tools solved the first half of that and stopped. You get a faithful account of every word, which is useful roughly never, because the question afterwards is not *what was said*. It is *what did I agree to, and what do I do first*. A 58-minute transcript answers neither.

The gap is not between talking and text. It is between text and the next action.

There is a second gap, during the call rather than after it. The moment you are least able to think clearly is the moment the room turns to you and asks what you think. Recap is reading the conversation anyway; it may as well be useful then, when help is worth something, rather than only in a summary email nobody opens.

The part most tools get wrong

An action item that reads "confirm the scope in writing by Friday" is a perfectly accurate record and a completely useless instruction. It names an outcome, not a beginning. For anyone who finds starting harder than finishing — and that is a lot of people, ADHD or otherwise — the distance between that sentence and the first keystroke is where the task quietly dies.

So Recap does not stop at the commitment. It asks how *you* work, and turns the commitment into a two-minute opening move: one concrete first step, a timebox, and something to do when you get stuck. Five work styles, chosen by the person, never inferred — because software guessing at how your attention works would be both wrong and unwelcome.

Hard to start

A first step small enough to be embarrassing, a timer, and permission to do it badly.

Needs a finish line

Define what done looks like first, then work backwards to the opening move.

One thing at a time

Split anything compound into its parts and name only the one that comes first.

Long uninterrupted block

Protect the time before the task, because the cost here is interruption rather than starting.


Think it through with someone

Name the person to put it in front of, and what to ask them for.

The same commitment produces a different plan under each one. That is the whole argument: the task did not change, and what it takes to begin it did.

One pass over a conversation

Recap listens to both sides of a call from your own machine, so nothing joins the meeting and no other participant sees a bot in the room.



Live coaching

When the room turns to you, a suggestion appears: how the moment reads, specific wording to use, why it works, and what to be careful of. Built from the last few things actually said, not a template.

Two-sided attribution

Your microphone and the call's audio are recorded separately, so "you" and "them" is a fact about where the sound came from rather than a diarisation guess.

Notes worth keeping

A headline, the decisions, action items with owners and timestamps, and the questions left open — which is usually the most valuable section.

Ask across every meeting

One question over everything you have recorded. Every answer cites the call and the second it came from, and an unsupported answer says so instead of guessing.

Recording and playback

The mixed audio is stored with the call and plays back against the transcript, so a disputed line can be listened to rather than argued about.

Shareable clips

Any moment can be cut to a link that opens without an account, exposing that clip and nothing else of the meeting around it.

Recordings, transcripts, summaries and workspace state are stored per account in PostgreSQL. Audio sits behind authentication and is never served from a shared cache. Passwords are scrypt hashes; the live audio slices sent for transcription are discarded, and only the finished recording is kept.

SCOPE

What is real, and what is not

Stating this plainly is part of the product, not a disclaimer bolted to it. The app itself says the same thing on the screen where it matters.

CAPABILITY	STATUS	DETAIL
Audio capture, both sides	REAL	Microphone plus shared tab or system audio, via the browser. Chrome or Edge on desktop.
Transcription	REAL	Whisper. Verbatim on ordinary speech; weaker on product jargon and proper nouns.
Live coaching	REAL	A model call on what has genuinely just been said in the room.
Summaries and action items	REAL	Generated from the stored transcript, with owners and deadlines attributed.
Work-style CTAs	REAL	Deterministic planner, with a model pass that asks a clarifying question when a commitment is vague.
Joining the meeting as a bot	NOT BUILT	Needs the Zoom Meeting SDK and marketplace approval. Recap listens from your machine instead.
Naming individual speakers	NOT BUILT	Without the platform's API there is no honest way to map a voice to a person. The far side is "Them".
Demo call	SCRIPTED	A scripted conversation that drives the real pipeline, so the product can be seen working without a second person.

ENGINEERING

Verified, not asserted

60

browser checks against the deployed app, zero console errors

38

server checks, including auth, rate limits and PII redaction

93

unit tests over the derived and resilience layers

The frontend is React and TypeScript on GitHub Pages. The backend is a dependency-free Node server on Heroku holding the API key, so no secret reaches the browser; the build fails if a key pattern appears in the bundle. Requests carry timeouts, retries with jittered backoff, a circuit breaker, idempotency keys and per-IP rate limits. Logs are structured and correlated, with transcript text and personal data dropped by denylist before anything is written.